

CTseg: A Tool for Brain CT Segmentation, Spatial Normalisation, and Volumetrics

Mikael Brudfors

NVIDIA

Abstract

This paper presents and validates CTseg, a freely available tool for brain CT segmentation, spatial normalisation, and volumetrics. CTseg builds on the Multi-Brain generative modelling framework, which introduces improved nonlinear registration, learned intensity priors, and groupwise atlas construction, to provide a CT-specific pipeline that produces tissue maps, deformation fields, and brain volume estimates in the same format as SPM’s unified segmentation, thereby extending SPM’s established analysis chain (voxel-based morphometry, atlas-based ROI analysis, lesion mapping) from MRI to CT. It is designed for routine hospital CT scans without requiring preprocessing or resampling in deployment. Although CTseg has been adopted in clinical research spanning, among other things, stroke, dementia, and brain morphometry, a systematic validation against an independent reference standard has been lacking. Using 59 paired MR/CT head scans from the SynthRAD2025 challenge, we evaluate CTseg across four dimensions: segmentation accuracy (in both native and normalised space), against an MRI-derived silver standard; spatial normalisation consistency through group-average sharpness and voxelwise coefficient of variation; brain volume agreement via intraclass correlation and Bland–Altman analysis; and downstream sex classification performance from normalised tissue maps. As a baseline, we apply SPM’s MRI-based unified segmentation directly to the CT images. CTseg significantly outperformed this baseline across segmentation, normalisation, and TBV agreement, while TIV agreement was comparable. CTseg is freely available at <https://github.com/WCHN/CTseg>, and all experiment code is included in the repository for full reproducibility.

1 Introduction

Computed tomography (CT) is the most frequently acquired neuroimaging modality in clinical practice. It is the first-line investigation for stroke, traumatic brain injury, and intracranial haemorrhage, and is routinely obtained in emergency, neurosurgical, and oncological settings. The vast majority of brain scans acquired worldwide are CT, not MRI, yet almost all computational neuroimaging tools have been developed for MRI.

Being able to automatically segment brain tissue and spatially normalise CT scans to a standard coordinate system [Friston et al., 1995] would unlock significant clinical and research value. It would enable population-level analyses such as voxel-based morphometry [Ashburner and Friston, 2000] on the large volumes of CT data that already exist in hospital archives, support quantitative assessment of brain atrophy from routine clinical scans, facilitate lesion mapping in standard space for stroke and trauma studies, and provide automated brain volume estimates (total brain volume, TBV, and total intracranial volume, TIV) that are recognised as clinical biomarkers. These capabilities are well established for MRI, where software suites such as SPM [Penny et al., 2007], FreeSurfer [Fischl, 2012], and FSL [Jenkinson et al., 2012] provide mature, validated pipelines. However, applying these MRI-trained tools to CT data is problematic because the tissue contrast in CT is fundamentally different: the Hounsfield unit scale reflects X-ray attenuation, yielding substantially poorer differentiation between grey matter (GM), white matter (WM), and cerebrospinal fluid (CSF) than MRI.

In this paper, we present and validate CTseg¹, a CT-specific tool that addresses these limitations. CTseg can be used within an SPM-compatible framework, without supervised labels or GPU dependence. Although CTseg has already been adopted in clinical research across diverse populations, a systematic validation against an independent reference standard has not yet been performed. We validate CTseg using 59 paired MR/CT head scans from the SynthRAD2025 challenge [Thummerer et al., 2025]. The paired design enables us to use SPM’s MRI segmentation [Ashburner and Friston, 2005] as a silver standard. As a baseline, we also apply SPM’s unified

¹Not to be confused with the unrelated “CTSeg” tool by [Adduru et al., 2020], which uses a different approach.

segmentation directly to the CT images, providing a controlled head-to-head comparison that isolates the benefit of CT-specific modelling. SPM is the appropriate primary comparator because both CTseg and SPM are unsupervised, model-based methods within the same framework. Deep learning methods for CT segmentation requires labels and generally address a narrower task: they produce segmentations but not the deformation fields or volumetric outputs needed for spatial normalisation, voxel-based morphometry, or brain volume estimation. A direct comparison with such methods is therefore outside the scope of this validation but would be valuable future work. CTseg also operates directly on the native scan geometry without resampling, avoiding partial volume effects that resampling to a fixed resolution can introduce.

2 Related Work

Automated analysis of brain CT has a shorter history than its MRI counterpart, largely because the poor soft-tissue contrast in CT was long considered insufficient for tissue-level segmentation. Early work focused primarily on spatial normalisation. Rorden et al. [2012] developed stroke-aged CT and MRI templates and packaged them in the SPM Clinical Toolbox, enabling affine and low-dimensional nonlinear normalisation of clinical CT scans for lesion mapping. Muschelli [2020] later produced a higher-resolution, unbiased CT template using an iterative registration procedure, together with coarse ventricle and whole-brain labels. Both approaches are template and normalisation resources rather than subject-level CT tissue-segmentation methods.

Tissue segmentation from CT has been addressed by a smaller body of work. Gupta et al. [2010] proposed an early approach for automatic GM, WM, and CSF segmentation from unenhanced CT using adaptive thresholding and anatomical knowledge, and Kemmling et al. [2012] introduced a probabilistic Hounsfield-unit decomposition method using MRI-derived tissue maps and a custom CT template. Manniesing et al. [2017] demonstrated that GM and WM can be segmented from 4D contrast-enhanced CT by exploiting temporal averaging across a perfusion series, followed by voxel feature extraction, SVM classification, and atlas-guided geodesic active contours. While effective for contrast-enhanced acquisitions, this method is not applicable to the far more common single-phase non-contrast CT.

More recently, deep learning methods have been applied to the problem. Srikrishna et al. [2021] trained a U-Net on paired CT/MRI data from the Gothenburg H70 Birth Cohort, using MRI-derived FreeSurfer labels as supervision, and achieved Dice scores of 0.79 (GM), 0.82 (WM), and 0.75 (CSF) on held-out CT images. Cai et al. [2020] used a deep-learning model to segment 11 intracranial structures from non-contrast head CT, targeting radiological reporting rather than tissue-class analysis. Gerken et al. [2023] trained a 2D U-Net for brain parenchyma and ventricular segmentation on CT, demonstrating robustness to pathology including haemorrhage and tumour. Son et al. [2024] used perceptual-loss U-Nets trained on 199 paired CT/MRI scans for GM, WM, and CSF segmentation. Most recently, Huisman et al. [2024] validated SynthSeg, a domain-randomised segmentation tool originally developed for MRI, on CT using 260 paired CT/MRI radiotherapy scans from five centres. These supervised and domain-randomised methods are segmentation-centric and do not themselves provide nonlinear deformation fields or an SPM-compatible end-to-end morphometry pipeline. Separately, Fielden et al. [2022] compared global CT-versus-MRI brain volumetrics using SPM and FSL in 69 paired scans, providing an independent reference point for CT-based volume estimation.

CTseg occupies a distinct position in this landscape: it is a generative model-based method that jointly estimates tissue segmentation and nonlinear spatial normalisation from a single non-contrast CT scan, without requiring labelled training data, contrast enhancement, or preprocessing. Its outputs are compatible with the standard SPM analysis pipeline, bridging the gap between CT acquisition and the quantitative analysis tools that are well established for MRI.

2.1 Prior Applications of CTseg

As listed in Table 3 of Appendix A, CTseg has been adopted in clinical research across diverse populations and applications, including stroke outcome prediction and perfusion mapping [Liu et al., 2021, Hoving et al., 2024, Luijten et al., 2026, Tangwiriyasakul et al., 2026], brain morphometry in Down syndrome [Sánchez-Moreno et al., 2024, 2025], COVID-19 neuroimaging [Duan et al., 2021], cardiac arrest prognosis [Gramespacher et al., 2024], radiomics in subarachnoid haemorrhage [Urbanos et al., 2025], PET imaging analysis [Tonietto et al., 2024], bone density estimation [Kalc et al., 2024], mortality prediction from routine hospital CT [Tsui et al., 2023], and dementia screening [Yelanchezian et al., 2025], collectively spanning over 5,000 subjects across 13 studies. These studies demonstrate the practical utility of CTseg across diverse clinical populations, but they do not substitute for direct validation against an independent reference standard. The present study provides such a validation.

3 CTseg

3.1 Overview

CTseg provides a CT-specific pipeline for simultaneous tissue segmentation and nonlinear spatial normalisation compatible with the SPM12 framework. It uses the the Multi-Brain (MB) generative modelling framework [Brudfors et al., 2020], which extends SPM’s unified segmentation approach [Ashburner and Friston, 2005] with improved affine+nonlinear registration and learned spatial (i.e. atlas) and intensity priors. Essentially, CTseg is an MB model trained on a mix of both MRI and CT brain scans. The inclusion of CT ensures the intensity priors capture Hounsfield-unit profiles, while MRI provides superior soft-tissue contrast for learning tissue boundaries. Full details of the learning the MB model used by CTseg are provided in Appendix B. The model produces six tissue class segmentations (GM, WM, CSF, bone, soft tissue, and background) together with nonlinear deformation fields that map each scan to a CT-derived atlas aligned with the MNI coordinate system. In addition, CTseg provides skull-stripping and estimates of TBV and TIV.

A key practical advantage is that CTseg works directly on clinical CT scans without any preprocessing. It handles arbitrary voxel sizes (including variable slice thickness common in clinical CT), arbitrary fields of view, and does not assume any particular orientation or prior alignment. The generative modelling approach also means that CTseg does not require manually labelled training data. Because CTseg produces outputs in the same format as SPM’s unified segmentation (native and template-space tissue maps, deformation fields, and modulated images), it serves as a drop-in replacement that extends SPM’s downstream analysis tools, including voxel-based morphometry [Ashburner and Friston, 2000], atlas-based region-of-interest analysis [Tzourio-Mazoyer et al., 2002], and voxel-based lesion–symptom mapping [Bates et al., 2003], from MRI to CT data.

4 Validation Study

4.1 Data

Validating a CT segmentation tool against an MRI-derived reference requires a dataset of spatially aligned, same-subject MR and CT brain scans. Such datasets are rare in the public domain, because MR and CT are typically acquired for different clinical indications and CT involves ionising radiation that is difficult to justify for research alone. The few existing paired datasets come from radiotherapy planning, where both modalities are acquired as part of routine care. We therefore used the SynthRAD2025 Grand Challenge dataset [Thummerer et al., 2025], which is, to our knowledge, among the largest publicly available paired MR–CT datasets. The SynthRAD2025 dataset comprises 2,362 paired MRI–CT and CBCT–CT cases from head-and-neck, thoracic, and abdominal cancer patients, acquired at five European university medical centres (UMC Groningen, UMC Utrecht, Radboud UMC, LMU University Hospital Munich, and University Hospital of Cologne) using diverse scanners and acquisition protocols. From this dataset, we used the Task 1 (Head & Neck) MRI–CT training subset. The full training set comprises 221 subjects from three centres (A, C, and D); however, centre D data (65 subjects) is distributed under a restricted licence that limits its use to the duration of the challenge, so we used the 156 freely available subjects from centres A ($n = 91$) and C ($n = 65$), each with a paired MR and CT scan.

The challenge organisers pre-processed all images prior to release [Thummerer et al., 2025]: (1) MR images were rigidly registered to the corresponding planning CT using Elastix; (2) facial structures were removed (defaced) for anonymisation; (3) images were resampled to $1 \times 1 \times 3$ mm voxel spacing; and (4) images were cropped to the patient outline. No bias field correction or intensity normalisation was applied to the MR images. Images were provided in MetaImage format (.mha) and stored as INT16. Notably, deformable registration of the CT to the MR, to correct for anatomical differences between the two scan sessions, was only performed by the challenge organisers for the validation and test sets, not for the training data. We therefore replicated this step ourselves (described in Section 4.2).

The accompanying metadata provides patient demographics and acquisition parameters. Of the 156 subjects, 64 were male, 27 female, and 65 had sex unlisted; age was available for 90 subjects (mean 64 years, range 29–87). CT scans were acquired on scanners from two manufacturers (Philips, $n = 118$; Siemens, $n = 38$) across four scanner models, at 120 kVp, with slice thicknesses of 1–3 mm. The dataset is publicly available at <https://synthrad2025.grand-challenge.org/> (Zenodo DOI: 10.5281/zenodo.14918089). Figure 1 shows example MR and CT views from a representative subject.

Example paired MR and CT — 1HNA042




Figure 1: Example paired MR (top) and CT (bottom) images from the SynthRAD2025 dataset (Task 1 Head & Neck, MRI-CT training subset), shown in axial, coronal, and sagittal views.

4.2 Preprocessing

All images were preprocessed using the `spm-hospital-preproc` toolbox². Preprocessing consisted of four steps: (1) rigid-body alignment of each image to MNI space; (2) co-registration of the MR and CT images within each subject pair using normalised mutual information, to ensure precise spatial correspondence; (3) cropping of both images to the SPM12 template field of view ($181 \times 217 \times 60$ voxels at the original $1 \times 1 \times 3$ mm voxel size); and (4) deformable registration of the CT to the MR to correct for residual anatomical differences between the two scan sessions (e.g. from weight change or organ motion), replicating the procedure used by the challenge organisers for the validation and test sets [Thummerer et al., 2025]. The deformable registration used the Elastix framework [Klein et al., 2010] via SimpleITK with a B-spline transform (final grid spacing 15 mm, three resolution levels), optimised with Mattes mutual information, a bending energy penalty, and a rigidity penalty, using the same parameter file³ as the challenge preprocessing pipeline. These steps ensure that the MR and CT images share the same voxel grid, so that tissue maps produced by different segmentation methods can be compared voxelwise. The resulting deformed CT was used as input for all subsequent CT segmentation steps. We note that this preprocessing was performed solely to enable the paired comparison in this validation study; CTseg itself does not require any such preprocessing when used in practice.

After preprocessing, of the 156 subjects from centres A and C, 97 were excluded during quality review because they had incomplete brain coverage (the head-and-neck radiotherapy field of view often truncates the vertex or skull base). The remaining 59 subjects were processed by all the three below pipelines, where all three methods succeeded for every subject, so no further exclusions were necessary.

4.3 Segmentation Pipelines

Three segmentation pipelines were applied to each subject:

SPM on MRI (silver standard). SPM12’s unified segmentation [Ashburner and Friston, 2005] was applied to each preprocessed MR image using default parameters, producing native-space probability maps for six tissue classes. We retained the GM, WM, and CSF maps as the silver-standard reference segmentations.

²<https://github.com/WTCN-computational-anatomy-group/spm-hospital-preproc>

³<https://github.com/SynthRAD2025/preprocessing>

SPM on CT (baseline). The same SPM12 pipeline was applied to each preprocessed CT image. Because SPM’s tissue priors were designed for MRI, this represents an out-of-domain application. Bias field regularisation was increased (`biasreg=10`) to suppress bias field estimation, since CT images lack the intensity inhomogeneity characteristic of MRI. The forward deformation field was saved and used to produce a spatially normalised CT image via SPM’s normalise-write module.

CTseg on CT. CTseg’s atlas (`mu_CTseg.nii`) was originally learned in a groupwise-optimal space that does not coincide with MNI. This is a consequence of the Multi-Brain framework, which learns the atlas jointly with the registration parameters and is therefore not constrained to any predefined coordinate system. While this groupwise space is optimal for the training population, it means that CTseg’s deformation fields do not map directly to MNI space. This poses two challenges: (1) for this validation study, comparing spatially normalised tissue maps produced by CTseg and SPM requires both to be in the same coordinate system; otherwise voxelwise overlap metrics are meaningless; and (2) more generally, many neuroimaging applications require outputs in MNI space for group analyses, ROI-based studies, and use of standard MNI atlases.

To address this, we nonlinearly registered the CTseg atlas into SPM atlas space by fitting the Multi-Brain model [Brudfors et al., 2020] to a MNI152 T1 1 mm atlas⁴ using uninformative intensity priors. The resulting deformation field was used to warp the CTseg atlas into SPM space at 1.5 mm isotropic resolution, producing a new atlas (`mu_CTseg_spm.nii`) that shares the same coordinate system as the SPM12 tissue probability maps. All subsequent CTseg analyses used this SPM-aligned atlas, ensuring that CTseg’s deformation fields map directly to SPM space without requiring an additional affine alignment step. Both the original groupwise atlas and the MNI-aligned version are provided with CTseg, allowing users to select the appropriate atlas for their application (Appendix B).

CTseg was applied to each preprocessed CT image using the SPM-aligned atlas with the output voxel size set to 1.5 mm (matching SPM’s atlas-space resolution). Native-space tissue maps and modulated warped tissue maps were retained, and the estimated deformation field was used to produce a spatially normalised CT image. All three pipelines also produced modulated warped GM, WM, and CSF tissue maps in MNI space at 1.5 mm resolution, used for the predictive validation experiment.

4.4 Evaluation

4.4.1 Segmentation accuracy

For each tissue class (GM, WM, CSF), we compared the CT-based segmentations against the MRI silver standard. Tissue probability maps were binarised at a threshold of 0.5 before computing overlap and surface distance metrics. We report three complementary measures:

- **Dice coefficient** [Dice, 1945]: volumetric overlap, ranging from 0 (none) to 1 (perfect).
- **95th percentile Hausdorff distance (HD95)** [Huttenlocher et al., 1993]: near-worst-case boundary error (mm), robust to outliers.
- **Average symmetric surface distance (ASSD)**: mean boundary displacement (mm).

4.4.2 Spatial normalisation quality

Normalisation quality was assessed in three ways. First, the group-average of all normalised CT images was computed separately for each deformation (SPM-MR, SPM-CT, and CTseg); a sharper average indicates more consistent normalisation. Second, we computed the voxelwise coefficient of variation ($\text{CoV} = \sigma/|\mu|$) across subjects within a brain mask (mean intensity > 20 HU); a lower mean CoV indicates more consistent normalisation. Third, we directly compared the spatial accuracy of the deformation fields by warping each subject’s native-space MR tissue maps (GM, WM, CSF) to template space using both the SPM-CT and CTseg deformations, then computing Dice and ASSD against the SPM-MR warped tissue maps (binarised at 0.5) as reference. This isolates the quality of the deformation from differences in the segmentation model, since all three methods warp the same native-space MR tissue maps.

⁴<https://github.com/Jfortin1/MNITemplate>

4.4.3 Brain volumetrics

TBV (GM + WM) and TIV (GM + WM + CSF) were computed for all three methods by summing native-space tissue probability maps weighted by voxel volume. Before summation, an intracranial mask was applied to exclude voxels outside the brain. The mask was derived from the CTseg atlas by computing the softmax tissue probabilities and thresholding the sum of GM, WM, and CSF at 0.5. This template-space mask was then warped to each subject’s native space using the SPM-MR deformation field. The same mask was applied to all three methods to ensure a fair comparison. This step is necessary because head CT scans typically extend well below the skull base, and without masking, tissue probability assigned to extracranial structures (e.g. cervical spine white matter, facial soft tissue) can inflate volume estimates, particularly for SPM, whose tissue priors were designed for brain MRI and may assign non-negligible probabilities to non-brain tissue when applied to CT. Agreement was assessed using ICC(3,1) (a two-way mixed-effects, single-measures intraclass correlation assessing consistency) [Shrout and Fleiss, 1979], Pearson correlation, and Bland–Altman analysis [Bland and Altman, 1986].

4.4.4 Predictive validation of normalisation

As an implicit measure of normalisation quality, we evaluated how well normalised tissue maps support sex classification, following the paradigm of [Brudfors et al., 2019]. The rationale is that better spatial normalisation should yield more anatomically consistent feature representations across subjects, improving discriminability in downstream analyses. CTseg’s diffeomorphic registration model is more expressive than SPM12’s default parameterisation, and its atlas was learned jointly from MRI and CT data; if these richer deformations produce more descriptive normalised tissue maps, this should be reflected in classification performance—potentially even outperforming SPM applied to MRI.

We used $n = 58$ subjects for the experiment (one of the 59 subjects lacked sex metadata and was excluded), Figure 2 shows their age and sex distributions. For each method, modulated warped GM, WM, and CSF maps (i.e. tissue probability maps warped to template space and scaled by the Jacobian determinant to preserve total tissue volume) were smoothed (8 mm FWHM), masked to intracranial voxels, and concatenated into a single feature vector per subject. These features were used directly (without dimensionality reduction) in a kernel-based Gaussian process (GP) classification framework implemented in PRoNTTo [Schrouff et al., 2013], which computes a linear kernel from the feature vectors. Mean centering was performed within each cross-validation fold to prevent information leakage. We used 10-fold cross-validation with identical folds across methods. Given the small sample, this analysis is exploratory: we report accuracy and AUC but do not perform formal statistical comparisons between methods. Age regression results using the same framework are reported in Appendix C.




Figure 2: Age and sex distributions of the 58 subjects retained for prediction analysis (i.e. those with complete tissue maps from all three pipelines and both age and sex metadata available).

Table 1: Segmentation accuracy compared to MRI silver standard. Values are median [IQR]. p -values from Wilcoxon signed-rank test (Bonferroni-corrected).

Metric	Method	GM	WM	CSF
Dice	SPM	0.531 [0.515, 0.561]	0.659 [0.647, 0.672]	0.367 [0.342, 0.405]
	CTseg	0.598 [0.582, 0.620]	0.708 [0.699, 0.726]	0.463 [0.436, 0.511]
	p -value	<0.001	<0.001	<0.001
HD95 (mm)	SPM	4.243 [4.123, 4.690]	7.810 [6.500, 8.602]	6.000 [5.385, 6.403]
	CTseg	3.162 [3.000, 3.162]	4.359 [4.123, 4.472]	5.099 [4.690, 6.403]
	p -value	<0.001	<0.001	0.171
ASSD (mm)	SPM	1.135 [1.072, 1.229]	1.818 [1.646, 1.953]	1.821 [1.704, 1.949]
	CTseg	0.850 [0.805, 0.916]	1.068 [1.018, 1.120]	1.496 [1.415, 1.656]
	p -value	<0.001	<0.001	<0.001

4.4.5 Runtime

Processing time was recorded for each method and subject (segmentation only, excluding I/O and warping).

4.5 Statistical Analysis

Paired Wilcoxon signed-rank tests (two-sided, Bonferroni-corrected for three tissue classes) compared segmentation metrics between SPM-on-CT and CTseg-on-CT; the same tests and correction procedure were applied to the normalisation metrics (warped-tissue Dice and ASSD). Segmentation and normalisation metrics are reported as median [interquartile range]. Volumetric agreement is reported as ICC with 95% confidence intervals, Pearson r , and Bland–Altman bias with 95% limits of agreement. Significance level: $\alpha = 0.05$.

5 Results

5.1 Segmentation Accuracy

Table 1 and Figure 3 summarise the segmentation metrics across all subjects and tissue classes. CTseg achieved significantly higher Dice scores and lower surface distances than SPM-on-CT for all three tissue classes ($p < 0.001$ for all Dice and ASSD comparisons; CSF HD95 did not reach significance). Figure 4 shows tissue probability maps for a representative subject, and Figure 5 shows segmentation overlays for the best and worst cases.

5.2 Spatial Normalisation

Figure 6 shows the group-average normalised CT images computed using the three deformation fields. Visual inspection suggests that the CTseg average preserves cortical detail and ventricular boundaries more clearly than the SPM-CT average, consistent with more consistent spatial normalisation across subjects. The mean voxelwise coefficient of variation (CoV) within the brain mask was 0.319 for CTseg versus 0.402 for SPM-CT, confirming more consistent alignment. (The SPM-MR CoV of 0.458 is not directly comparable because MRI has a wider dynamic range than CT, inflating voxelwise variability.) To quantify normalisation accuracy independently of segmentation differences, we warped each subject’s native-space MR tissue maps to template space using both the SPM-CT and CTseg deformations, then compared against the SPM-MR warped maps as reference (Figure 7). CTseg achieved significantly higher Dice overlap than SPM-CT for GM (0.660 vs. 0.633, $p < 0.001$) and WM (0.796 vs. 0.764, $p < 0.001$), with a smaller but significant improvement for CSF (0.642 vs. 0.636, $p = 0.031$). ASSD was significantly lower for CTseg in GM (0.903 vs. 0.970 mm) and WM (0.855 vs. 0.932 mm), with no significant difference for CSF. These results confirm that CTseg’s deformation fields more accurately align brain anatomy to template space than those obtained by applying SPM’s MRI-trained model to CT.

5.3 Brain Volumetrics

Table 2 and Figure 8 report volumetric agreement between CTseg and MRI-derived estimates. For TBV, CTseg achieves a substantially higher ICC (0.838 vs. 0.653) and tighter limits of agreement than SPM-CT, indicating that it better preserves the relative ordering of brain sizes across subjects. CTseg shows a systematic negative bias




Figure 3: Segmentation metrics comparing SPM and CTseg on CT against the MRI silver standard. Left: Dice coefficient (higher is better). Centre: HD95 (lower is better). Right: ASSD (lower is better).

Table 2: Agreement of brain volume estimates (CT methods vs. MRI reference).

Measure	Method	ICC [95% CI]	r	Bias (ml)	LoA (ml)
TBV	SPM	0.653 [0.478, 0.778]	0.672	3.2	[-153.1, 159.5]
TBV	CTseg	0.838 [0.741, 0.900]	0.838	-53.8	[-149.0, 41.4]
TIV	SPM	0.804 [0.691, 0.878]	0.813	-196.5	[-329.2, -63.8]
TIV	CTseg	0.826 [0.724, 0.893]	0.828	-227.6	[-347.1, -108.0]

of approximately -54 ml, whereas SPM-CT’s bias is near zero (3 ml). For TIV, both methods underestimate the MR-derived reference by approximately 197–228 ml, with similar ICC values (SPM-CT: 0.804, CTseg: 0.826). The TIV underestimation is expected given the inherently poor contrast between CSF and brain parenchyma on CT.

5.4 Predictive Validation

Figure 9 shows the sex classification ROC curves. CTseg achieved the highest AUC (0.924), outperforming both SPM-on-MR (0.886) and SPM-on-CT (0.869). This suggests that CTseg’s more expressive diffeomorphic registration model preserves sex-discriminative morphological features (e.g. total brain size, cortical thickness patterns) more faithfully than either SPM pipeline. Notably, CTseg applied to CT even outperformed SPM applied to the higher-contrast MRI, consistent with the hypothesis that the richer deformation model yields more descriptive normalised tissue maps. Importantly, AUC is threshold-independent and robust to the class imbalance in our dataset (70% male, 30% female): a classifier that simply predicts the majority class would achieve 70% accuracy but an AUC of only 0.5.

5.5 Runtime

All experiments were run on a workstation with an Intel Core Ultra 7 165H (3.80 GHz) processor and 32 GB RAM, running Windows 11. No GPU acceleration was used. Median segmentation time per subject was 84 s (mean




Figure 4: Tissue probability maps for a representative subject (near-median Dice). Rows: GM, WM, CSF. Columns: SPM on MRI (silver standard), SPM on CT, CTseg on CT. Looking at, e.g., GM of CTseg, it clearly shows finer details in the cortex compared to SPM on CT.




Figure 5: Segmentation overlays for best and worst cases (axial view). Rows: SPM on MRI (silver standard), SPM on CT, CTseg on CT. Columns: best and worst by SPM mean Dice (left pair), best and worst by CTseg mean Dice (right pair). Tissue colours: GM (red), WM (blue), CSF (green). Mean Dice scores are shown in column headers.




Figure 6: Group-average normalised CT (axial view) for each deformation method. Sharper anatomical boundaries indicate more consistent spatial normalisation across subjects. Red boxes highlight the lateral ventricle and insular cortex region, where the CTseg average preserves cortical sulcal detail and the grey-white matter boundary comparably to the SPM-MR reference, whereas the SPM-CT average is visibly more blurred.

83 $\pm$ 11 s) for SPM-MR, 100 s (mean 103 $\pm$ 17 s) for SPM-CT, and 163 s (mean 173 $\pm$ 36 s) for CTseg ($n = 59$). CTseg is approximately twice as slow as SPM, which is expected given its more flexible nonlinear registration model (Multi-Brain diffeomorphic framework versus SPM’s DCT-based normalisation). All three methods complete within minutes on a standard workstation without GPU acceleration, making them practical for both research and clinical deployment.

6 Discussion

CTseg addresses a different problem from supervised CT labelling tools: it provides a CT-specific generative pipeline for tissue segmentation, spatial normalisation, and brain volumetrics within an SPM-compatible framework, bridging the gap between the large volume of clinical CT data and the quantitative analysis tools established for MRI. The goal of this validation is to demonstrate that CTseg produces reliable outputs across four complementary evaluation dimensions, substantially outperforming the natural baseline of applying SPM’s MRI-trained model to CT. Beyond this formal validation, CTseg’s practical utility is supported by its adoption in over a dozen clinical studies spanning stroke, dementia, brain morphometry, and other domains (Section 2.1), where CTseg-derived outputs have produced statistically significant and clinically interpretable findings.

The results largely confirm this. CTseg showed improvements over SPM-on-CT in most comparisons: significantly higher Dice scores for all tissue classes, more accurate spatial normalisation by warped-tissue overlap, and higher TBV agreement. However, CSF surface distances were not uniformly improved (HD95 did not differ significantly), TIV agreement was comparable between methods, and the predictive validation, while favouring CTseg, is limited by the small sample size.

For segmentation, CTseg achieved significantly higher Dice scores for all three tissue classes (GM: 0.598 vs. 0.531; WM: 0.708 vs. 0.659; CSF: 0.463 vs. 0.367; all $p < 0.001$), with correspondingly lower surface distances. These scores are lower than those reported for supervised deep learning methods trained on paired CT/MRI data [Srikrishna et al., 2021], which is expected: CTseg is an unsupervised generative model that requires no labelled training data and jointly estimates segmentation, registration, and bias field parameters. The trade-off is generality: CTseg is designed to work directly on routine non-contrast CT scans, whereas supervised methods require retraining for new populations, scanners, or acquisition protocols. The largest absolute improvement over SPM-CT was in CSF, where the poor CT contrast between CSF and parenchyma poses the greatest challenge for MRI-trained models. That even GM and WM, which are better differentiated on CT than CSF, show significant improvement underscores the value of CT-specific intensity priors and atlas geometry.

Spatial normalisation quality showed a similar pattern: CTseg produced sharper group averages and higher Dice




Figure 7: Normalisation metrics comparing SPM-CT and CTseg deformations. Native-space MR tissue maps were warped to template space using each deformation and compared against the SPM-MR warped tissue maps as reference. Left: Dice coefficient (higher is better). Right: ASSD (lower is better). HD95 is omitted because at the 1.5 mm template resolution, the distance transform quantises surface distances into a few discrete values, making HD95 insensitive to differences between methods.




Figure 8: Bland–Altman plots for TBV (left) and TIV (right). Solid line: mean bias. Dashed lines: 95% limits of agreement.




Figure 9: Sex classification ROC curves. Higher AUC indicates better normalisation quality.

overlap when the same native-space MR tissue maps were warped using each method’s deformation field. This confirms that the improvement is not merely due to different segmentation models but reflects genuinely more accurate anatomical alignment.

For brain volumetrics, CTseg achieved a substantially higher ICC for TBV (0.838 vs. 0.653), indicating that it better preserves the relative ordering of brain sizes across subjects. CTseg’s systematic negative bias (−54 ml for TBV, −228 ml for TIV) is partially attributable to its atlas class definitions and the inherently poor CSF contrast on CT (discussed further in Limitations). The predictive validation provides a complementary application-level assessment: CTseg achieved the highest sex classification AUC (0.924 vs. 0.869 for SPM-CT and 0.886 for SPM-MR), suggesting that its more expressive registration model produces normalised tissue maps that better preserve morphological features relevant to downstream analyses, even outperforming SPM applied to the higher-contrast MRI data.

6.1 Limitations

Several limitations should be acknowledged.

Reference standard and evaluation scope. Our MRI-derived reference segmentations constitute a *silver* standard, not a gold standard; both CT-based methods may be penalised for disagreeing with the MRI reference where it is itself inaccurate. We validate only GM, WM, and CSF, the classes for which MRI provides a credible reference, and skull stripping, although one of CTseg’s outputs, was not directly evaluated. We also do not compare against deep learning-based methods: as discussed in Section 1, no established deep learning tool offers the same joint segmentation, normalisation, and volumetric pipeline, and a comparison between supervised and unsupervised paradigms would be better addressed in dedicated future work.

Dataset and sample size. The SynthRAD2025 dataset consists of head-and-neck radiotherapy planning scans from cancer patients, not a general neuroimaging cohort. Publicly available paired MR–CT brain datasets are

exceedingly rare, so co-registered same-subject data almost exclusively comes from radiotherapy workflows; SynthRAD2025 is, to our knowledge, among the largest such resources for the head region. Of the 156 available subjects, 97 (62%) were excluded for incomplete brain coverage (a consequence of the radiotherapy planning field of view, which often truncates the vertex or skull base), leaving 59 for quantitative analysis. This limits statistical power and generalisability, although the successful application of CTseg across diverse clinical populations (Appendix A) provides indirect evidence of broader applicability. The predictive validation is limited by the small sample size ($n = 58$) and the imbalanced sex distribution (approximately 70% male) in this retained cohort.

Validation setting versus deployment. The validation involved substantial preprocessing (rigid alignment to MNI, MR-CT co-registration, deformable registration of CT to MR, and resampling to 1 mm isotropic), whereas CTseg is intended for direct application to unprocessed clinical CT. Native-space clinical robustness is therefore not directly tested here. Additionally, the validation uses the MNI-aligned 1.5 mm atlas configuration, not CTseg’s original groupwise atlas; results may differ with the original atlas, particularly for scans with extended inferior coverage. Furthermore, the nonlinear alignment of the CTseg atlas to MNI space was performed by fitting to a T1-weighted MRI atlas. Although the original CTseg atlas was learned from both MRI and CT data [Brudfors et al., 2020] and is not inherently biased toward either modality, aligning it to an MR-based MNI atlas may introduce a small MR-specific bias in the atlas shape. In practice, this bias is likely negligible, as the primary anatomical structures are consistent across modalities, but it should be noted. CTseg provides both the original groupwise atlas and the MNI-aligned version.

Volumetric interpretation. If SPM-CT shows closer volumetric agreement with MRI-derived volumes than CTseg does, this should be interpreted with caution: the MRI reference volumes are themselves produced by SPM’s unified segmentation, which shares the same tissue probability atlas, Gaussian mixture structure, regularisation, and cleanup pipeline as SPM-CT. Two runs of the same generative model with the same priors will produce correlated systematic biases, even when the input modality differs. Closer volume agreement between SPM-CT and SPM-MR may therefore reflect shared-method bias rather than genuinely superior accuracy. CTseg, by contrast, uses a different atlas, different priors, and a different registration framework (Multi-Brain), so its volume estimates are more independent of the MRI-derived reference. The observed systematic bias suggests that post hoc calibration may be feasible, but this was not evaluated here; by contrast, the lower correlation of SPM-CT reflects poorer rank-order agreement that cannot be corrected by a linear offset. CT also provides substantially lower contrast between CSF and brain parenchyma than MRI, which may cause CTseg to under-segment CSF. Since TIV is defined as $GM + WM + CSF$, any CSF under-segmentation directly reduces the TIV estimate; TBV may also be affected if partial-volume voxels at tissue boundaries are misclassified. This is an inherent limitation of CT-based brain segmentation, not specific to CTseg’s algorithm. A further source of volume discrepancy specific to CTseg is its atlas tissue class definitions: structures such as ventricular calcifications (commonly found in the choroid plexus) and the venous sinuses are modelled as non-brain classes (bone or soft tissue rather than CSF or GM/WM), as visible in the CTseg atlas (Figure 12). Because these structures are intracranial but assigned to non-brain tissue classes, they are excluded from TBV and TIV estimates, contributing to the systematic negative bias relative to the MRI reference.

6.2 Future Work

Valuable extensions include multi-site validation, comparison with deep learning approaches, validation of bone/soft-tissue classes, and evaluation on pathological populations where robust CT segmentation is most needed. Furthermore, CTseg’s generative framework is not inherently specific to CT: the tissue probability atlas and spatial normalisation machinery are modality-agnostic, and only the intensity prior is CT-specific. By training intensity priors on T1-weighted (or other MR contrast) images, the same framework could be applied directly to MRI data, offering a unified generative segmentation and normalisation tool across imaging modalities.

7 Conclusion

We have validated CTseg, a CT-specific brain segmentation and spatial normalisation tool, using 59 paired MR/CT scans from the SynthRAD2025 challenge. Relative to SPM’s MRI-based unified segmentation applied to CT, CTseg achieved significantly higher Dice scores for all tissue classes (GM, WM, CSF), more accurate spatial normalisation as measured by warped-tissue overlap, and stronger TBV agreement with the MRI reference ($ICC = 0.838$ vs. 0.653). TIV agreement was comparable between methods, with both showing substantial negative bias attributable to the

inherently poor CSF contrast on CT. Sex classification from normalised tissue maps favoured CTseg (AUC = 0.924), though the small sample limits definitive conclusions. These results, combined with CTseg’s growing adoption in clinical research (Appendix A), support its use for CT brain morphometry in settings similar to those evaluated here.

References

- V Adduru, S A Baum, C Zhang, M Helguera, R Zand, M Lichtenstein, C J Griessenauer, and A M Michael. A method to estimate brain volume from head CT images and application to detect brain atrophy in Alzheimer disease. *American Journal of Neuroradiology*, 41(2):224–230, 2020. doi: 10.3174/ajnr.A6402.
- John Ashburner and Karl J Friston. Voxel-based morphometry—the methods. *NeuroImage*, 11(6):805–821, 2000.
- John Ashburner and Karl J Friston. Unified segmentation. *NeuroImage*, 26(3):839–851, 2005.
- Elizabeth Bates, Stephen M Wilson, Ayse Pinar Saygin, Frederic Dick, Martin I Sereno, Robert T Knight, and Nina F Dronkers. Voxel-based lesion–symptom mapping. *Nature Neuroscience*, 6(5):448–450, 2003. doi: 10.1038/nm1050.
- J Martin Bland and Douglas G Altman. Statistical methods for assessing agreement between two methods of clinical measurement. *The Lancet*, 327(8476):307–310, 1986.
- Mikael Brudfors. *Generative Models for Preprocessing of Hospital Brain Scans*. PhD thesis, UCL (University College London), 2020.
- Mikael Brudfors, Yaël Balbastre, Parashkev Nachev, and John Ashburner. A tool for super-resolving multimodal clinical MRI. *arXiv preprint arXiv:1909.01140*, 2019.
- Mikael Brudfors, Yaël Balbastre, Guillaume Flandin, Parashkev Nachev, and John Ashburner. Flexible Bayesian modelling for nonlinear image registration. In *International Conference on Medical Image Computing and Computer Assisted Intervention*, pages 253–263. Springer, 2020.
- Jason C Cai, Zeynettin Akkus, Kenneth A Philbrick, Arunni Boonrod, Safa Hoodeshenas, Alexander D Weston, Pouria Rouzrokh, Gian Marco Conte, Atefeh Zeinoddini, David C Vogelsang, Qiao Huang, and Bradley J Erickson. Fully automated segmentation of head CT neuroanatomy using deep learning. *Radiology: Artificial Intelligence*, 2(5):e190183, 2020. doi: 10.1148/ryai.2020190183.
- Sasank Chilamkurthy et al. Deep learning algorithms for detection of critical findings in head CT scans: a retrospective study. *The Lancet*, 392(10162):2388–2396, 2018.
- Lee R Dice. Measures of the amount of ecologic association between species. *Ecology*, 26(3):297–302, 1945.
- Kuaikuai Duan, Enrico Premi, Andrea Pilotto, Viviana Cristillo, Alberto Benussi, Ilenia Libri, Marcello Giunta, H. Jeremy Bockholt, Jingyu Liu, Riccardo Campora, Alessandro Pezzini, Roberto Gasparotti, Mauro Magoni, Alessandro Padovani, and Vince D. Calhoun. Alterations of frontal-temporal gray matter volume associate with clinical measures of older adults with COVID-19. *Neurobiology of Stress*, 14:100326, 2021. doi: 10.1016/j.ynstr.2021.100326.
- SW Fielden, D Beiler, KA Cauley, and V Troiani. A comparison of global brain volumetrics obtained from CT versus MRI using 2 publicly available software packages. *American Journal of Neuroradiology*, 43(2):245–250, 2022. doi: 10.3174/ajnr.A7403.
- Bruce Fischl. FreeSurfer. *NeuroImage*, 62(2):774–781, 2012.
- Karl J Friston, John Ashburner, Christopher D Frith, Jean-Baptiste Poline, John D Heather, and Richard SJ Frackowiak. Spatial registration and normalization of images. *Human Brain Mapping*, 3(3):165–189, 1995. doi: 10.1002/hbm.460030303.
- Annika Gerken, Sina Walluscheck, Peter Kohlmann, Ivana Galinovic, Kersten Villringer, Jochen B Fiebach, Jan Klein, and Stefan Heldmann. Deep learning-based segmentation of brain parenchyma and ventricular system in CT scans in the presence of anomalies. *Frontiers in Neuroimaging*, 2:1228255, 2023. doi: 10.3389/fnimg.2023.1228255.

- Hannes Gramespacher, Maximilian HT Schmieschek, Clemens Warnke, Christoph Adler, et al. Analysis of cerebral CT based on supervised machine learning as a predictor of outcome after out-of-hospital cardiac arrest. *Neurology*, 103(1):e209583, 2024. doi: 10.1212/WNL.0000000000209583.
- Varsha Gupta, Wojciech Ambrosius, Guoyu Qian, Anna Blazejewska, Radoslaw Kazmierski, Andrzej Urbanik, and Wieslaw L Nowinski. Automatic segmentation of cerebrospinal fluid, white and gray matter in unenhanced computed tomography images. *Academic Radiology*, 17(11):1350–1358, 2010. doi: 10.1016/j.acra.2010.06.005.
- Jan W Hoving, Praneeta R Konduri, Manon L Tolhuisen, et al. Impact of intracranial volume and brain volume on the prognostic value of computed tomography perfusion core volume in acute ischemic stroke. *Journal of Cardiovascular Development and Disease*, 11(3):80, 2024.
- Selena Huisman, Matteo Maspero, Marielle Philippens, Joost Verhoeff, and Szabolcs David. Validation of SynthSeg segmentation performance on CT using paired MRI from radiotherapy patients. *NeuroImage*, 303:120922, 2024. doi: 10.1016/j.neuroimage.2024.120922.
- Daniel P Huttenlocher, Gregory A Klanderman, and William J Rucklidge. Comparing images using the Hausdorff distance. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 15(9):850–863, 1993.
- IXI. IXI dataset. <http://brain-development.org/ixi-dataset/>, 2023.
- Mark Jenkinson, Christian F Beckmann, Timothy EJ Behrens, Mark W Woolrich, and Stephen M Smith. FSL. *NeuroImage*, 62(2):782–790, 2012.
- Polona Kalc, Felix Hoffstaedter, Eileen Luders, Christian Gaser, and Robert Dahnke. Approximation of bone mineral density and subcutaneous adiposity using T1-weighted images of the human head. *Imaging Neuroscience*, 2:1–22, 2024. doi: 10.1162/imag_a_00390.
- A Kemmling, H Wersching, K Berger, S Knecht, C Groden, and I Nölte. Decomposing the Hounsfield unit: probabilistic segmentation of brain tissue in computed tomography. *Clinical Neuroradiology*, 22(1):79–91, 2012. doi: 10.1007/s00062-011-0123-0.
- Stefan Klein, Marius Staring, Keelin Murphy, Max A Viergever, and Josien PW Pluim. Elastix: a toolbox for intensity-based medical image registration. *IEEE Transactions on Medical Imaging*, 29(1):196–205, 2010.
- Hugo J Kuijf, Edwin Bennink, Koen L Vincken, Nick Weaver, Geert Jan Biessels, and Max A Viergever. MR brain segmentation challenge 2018 data, 2024. Version 1.0.
- Bennett Landman and Simon Warfield. MICCAI 2012 workshop on multi-atlas labeling. In *Proc. MICCAI Grand Challenge on Multi-Atlas Labeling*, 2012.
- Chang Liu, Hansheng Liu, Deping Wu, Zhiming Zhou, WenGuo Huang, Zhilin Wu, Wenjie Zi, and Qingwu Yang. Severe brain atrophy predicts poor clinical outcome after endovascular treatment of acute basilar artery occlusion. *Frontiers in Aging Neuroscience*, 13:720061, 2021. doi: 10.3389/fnagi.2021.720061.
- Sven P R Luijten, Aravind Ganesh, Adam P Marcus, et al. A CT imaging-based prediction model of functional outcome and benefit of endovascular thrombectomy for ischemic stroke. *European Radiology*, 2026. doi: 10.1007/s00330-025-12207-7.
- Rashindra Manniesing, Marcel T H Oei, Luuk J Oostveen, Jaime Melendez, Ewoud J Smit, Bram Platel, Clara I Sánchez, Frederick J A Meijer, Mathias Prokop, and Bram van Ginneken. White matter and gray matter segmentation in 4D computed tomography. *Scientific Reports*, 7:119, 2017.
- John Muschelli. A publicly available, high resolution, unbiased CT brain template. In *International Conference on Information Processing and Management of Uncertainty in Knowledge-Based Systems*, pages 358–366. Springer, 2020.
- William D Penny, Karl J Friston, John T Ashburner, Stefan J Kiebel, and Thomas E Nichols. *Statistical Parametric Mapping: The Analysis of Functional Brain Images*. Academic Press, 2007.
- Christopher Rorden, Leonardo Bonilha, Julius Fridriksson, Benjamin Bender, and Hans-Otto Karnath. Age-specific CT and MRI templates for spatial normalization. *NeuroImage*, 61(4):957–965, 2012.

- Beatriz Sánchez-Moreno, Linda Zhang, Gloria Mateo, et al. Voxel-based dysconnectomic brain morphometry with computed tomography in Down syndrome. *Annals of Clinical and Translational Neurology*, 11(1):143–155, 2024. doi: 10.1002/acn3.51940.
- Beatriz Sánchez-Moreno, Linda Zhang, Ana Carril Salaberry, et al. Performance of computed tomography and magnetic resonance morphometry in evaluating brain atrophy in Down syndrome. *Alzheimer’s & Dementia*, 21:e70296, 2025. doi: 10.1002/alz.70296.
- Jessica Schrouff, Maria J Rosa, Jane M Rondina, Andre F Marquand, Carlton Chu, John Ashburner, Christophe Phillips, Jonas Richiardi, and Janaina Mourão-Miranda. PRoNTo: pattern recognition for neuroimaging toolbox. *Neuroinformatics*, 11(3):319–337, 2013.
- Patrick E Shrout and Joseph L Fleiss. Intraclass correlations: uses in assessing rater reliability. *Psychological Bulletin*, 86(2):420–428, 1979.
- Won Jun Son, Sung Jun Ahn, Ji Young Lee, and Hyunyeol Lee. Automated brain segmentation on computed tomographic images using perceptual loss based convolutional neural networks. *Investigative Magnetic Resonance Imaging*, 28(4):193–201, 2024. doi: 10.13104/imri.2024.0023.
- Meera Srikrishna, Joana B Pereira, Rolf A Heckemann, Giovanni Volpe, Danielle van Westen, Anna Zettergren, Silke Kern, Lars-Olof Wahlund, Eric Westman, Ingmar Skoog, and Michael Schöll. Deep learning from MRI-derived labels enables automatic brain tissue classification on human brain CT. *NeuroImage*, 244:118606, 2021.
- Chayanin Tangwiriyaakul, Pedro Borges, Guilherme Pombo, Stefano Moriconi, Michael S Elmaleh, et al. Deep generative computed perfusion-deficit mapping of ischaemic stroke. *Communications Biology*, 9(1):219, 2026. doi: 10.1038/s42003-025-09495-6.
- Adrian Thummerer, Erik van der Bijl, Arthur Jr Galapon, Florian Kamp, Mark Savenije, Christina Muijs, Shafak Aluwini, Roel J H M Steenbakkens, Stephanie Beuel, Martijn P W Intven, Johannes A Langendijk, Stefan Both, Stefanie Corradini, Viktor Rogowski, Maarten Terpstra, Niklas Wahl, Christopher Kurz, Guillaume Landry, and Matteo Maspero. SynthRAD2025 grand challenge dataset: Generating synthetic CTs for radiotherapy from head to abdomen. *Medical Physics*, 52(7):e17981, 2025. doi: 10.1002/mp.17981.
- Matteo Tonietto, Oscar Sotolongo-Grau, Núria Roé-Vellvé, Santiago Bullich, Juan Pablo Tartari, Ángela Sanabria, Ainhoa García-Sánchez, Edilio Borroni, Christopher Galli, Esther Pérez-Martínez, Joan Castell-Conesa, Isabel Roca, Lluís Tárraga, Agustín Ruiz, Andrew W. Stephens, Mercè Boada, Gregory Klein, and Marta Marquié. Head-to-head comparison of tau PET tracers [18F]PI-2620 and [18F]RO948 in non-demented individuals with brain amyloid deposition: the TAU-PET FACEHBI cohort. *Alzheimer’s Research & Therapy*, 16(1):257, 2024. doi: 10.1186/s13195-024-01622-5.
- Alex Tsui, Petru-Daniel Tudosi, Mikael Brudfors, et al. Predicting mortality in acutely hospitalised older patients: the impact of model dimensionality. *BMC Medicine*, 21:10, 2023. doi: 10.1186/s12916-022-02698-2.
- Nathalie Tzourio-Mazoyer, Brigitte Landeau, Dimitri Papathanassiou, Fabrice Crivello, Olivier Etard, Nicolas Delcroix, Bernard Mazoyer, and Marc Joliot. Automated anatomical labeling of activations in SPM using a macroscopic anatomical parcellation of the MNI MRI single-subject brain. *NeuroImage*, 15(1):273–289, 2002. doi: 10.1006/nimg.2001.0978.
- Gemma Urbanos, Ana M Castaño-León, et al. Comprehensive predictive modeling in subarachnoid hemorrhage: integrating radiomics and clinical variables. *Neurosurgical Review*, 2025. doi: 10.1007/s10143-025-03679-8.
- Mukish Yelanchezian, Cristian Gonzalez-Prieto, et al. The diagnostic accuracy of CTseg segmentation software for dementia in a New Zealand memory service. *Journal of Alzheimer’s Disease Reports*, 9:25424823251332448, 2025. doi: 10.1177/25424823251332448.

A Published Studies Using CTseg

Table 3 lists published studies that have used CTseg or CTseg-derived segmentations, spanning brain morphometry, stroke outcome prediction, dementia screening, cardiac arrest prognosis, radiomics, PET imaging, bone density estimation, and mortality prediction, collectively encompassing over 5,000 subjects.

Table 3: Published external applications of CTseg.

Study	Domain	CTseg Use	N	Key Finding
Liu et al. [2021]	Basilar artery occlusion	BPV, CSF volume, brain atrophy index	231	Severe brain atrophy (estimated from CT via CTseg) independently predicted poor functional outcome after endovascular treatment in the BASILAR registry.
Hoving et al. [2024]	Ischaemic stroke	ICV, TBV from baseline NCCT	200	Adjusting CT perfusion core volume by CTseg-derived ICV or TBV did not improve prognostic value beyond absolute core volume in the MR CLEAN Registry.
Luijten et al. [2026]	Ischaemic stroke (EVT)	Brain parenchymal fraction	1,391	CT-based prediction model incorporating CTseg-derived brain atrophy improved identification of patients benefiting from thrombectomy across 7 randomised trials.
Sánchez-Moreno et al. [2024]	Down syndrome	VBM, tissue volumes, spatial normalisation	98	First CT-based VBM study in Down syndrome; CTseg-derived GM volumes correlated with dementia severity and plasma AD biomarkers (p-tau181, NfL) in hippocampal and cingulate regions.
Sánchez-Moreno et al. [2025]	Down syndrome	CT vs MRI morphometry comparison	23	92.3% agreement between CTseg (CT) and SPM DARTEL (MRI) morphometry; similar patterns of regional volume loss associated with cognitive decline.
Urbanos et al. [2025]	Subarachnoid haemorrhage	GM/WM segmentation for radiomics feature extraction	403	Radiomics features extracted from CTseg tissue segmentations predicted 6-month mortality, functional outcome (GOS), vasospasm, and hydrocephalus in SAH patients.
Kalc et al. [2024]	Bone mineral density	CT skull segmentation for validation	114	CTseg skull segmentation of OASIS-3 CT scans used as ground truth to validate MRI-derived proxy measures of bone mineral density and subcutaneous adiposity.
Tsui et al. [2023]	Acute hospital mortality	TBV, TIV from routine NCCT	804	CT-derived brain volumes integrated into a multidimensional clinical model predicted mortality in acutely hospitalised older patients, demonstrating prognostic value of routinely acquired CT data.
Gramespacher et al. [2024]	Cardiac arrest outcome	GM segmentation from NCCT	132	CTseg GM segmentation of cerebral CT used as input to a supervised ML classifier predicting neurological outcome after out-of-hospital cardiac arrest.

Continued on next page

Study	Domain	CTseg Use	N	Key Finding
Tangwiriyasakul et al. [2026]	Ischaemic stroke perfusion	Spatial normalisation, GM/WM atlases	1,393	CTseg used to co-register and nonlinearly normalise routine NCCT and CTA to MNI space; CTseg GM/WM atlases used to segment derived perfusion-deficit substrates in a deep generative model.
Duan et al. [2021]	COVID-19 neuroimaging	VBM, tissue volumes, spatial normalisation	120	CT images normalised to MNI and segmented into six tissue classes using CTseg; VBM revealed frontal-temporal GM volume reductions associated with clinical severity in older COVID-19 patients.
Tonietto et al. [2024]	Alzheimer’s tau PET imaging	Skull/meninges segmentation from CT	18	CTseg skull probability maps from PET attenuation-correction CT used to define skull/meninges ROI for isolating off-target tau tracer binding in a head-to-head comparison of [^{18}F]PI-2620 and [^{18}F]RO948.
Yelanchezian et al. [2025]	Dementia screening	TBV, hippocampal volume	168	CTseg-derived TBV and hippocampal volume from routine clinical CT differentiated dementia from non-dementia (combined AUC = 0.74) in a New Zealand memory service.
Total				5,095

B Atlas Construction

The CTseg atlas and intensity priors were learned jointly from both CT and MRI data using the Multi-Brain framework [Brudfors et al., 2020]⁵. The training data comprised four publicly available datasets:

- **CQ500** [Chilamkurthy et al., 2018]: 491 non-contrast head CT scans, of which only the healthy subjects ($n \approx 54$) were used for atlas construction. In-plane resolution ~ 0.5 mm, variable slice thickness (0.625–5.0 mm).
- **IXI** [IXI, 2023]: 577 multi-sequence MR images (T1w, T2w, PDw) from healthy subjects, acquired at three London hospitals, ~ 1 mm isotropic.
- **MICCAI 2012 Multi-Atlas Labeling** [Landman and Warfield, 2012]: 35 T1-weighted MRI scans (30 unique subjects, 5 rescanned) with expert manual segmentations of 133 brain regions, 1 mm isotropic. The manual labels were combined into GM and WM tissue classes and used to supervise the atlas learning.
- **MRBrainS18** [Kuijf et al., 2024]: 7 training subjects with multi-sequence MRI (T1w, T1-IR, T2-FLAIR) and manual segmentations of 8 brain structures, from subjects aged over 50, including patients with pathologies. The manual labels were combined into GM (cortical grey matter + basal ganglia), WM (white matter + white matter lesions), and CSF (cerebrospinal fluid + ventricles) to supervise the atlas learning. Pathological tissue (e.g. infarctions) was masked out during training.

The inclusion of both CT and MRI in the training set is a key design choice: MRI provides superior soft-tissue contrast for learning tissue boundaries, while CT improves learning of bone structures and ensures the intensity priors capture Hounsfield-unit profiles. Where manual segmentation labels were available (MICCAI 2012, MRBrainS18), they were used to guide the merging of learned Gaussian components into tissue classes; the remaining datasets (CQ500, IXI) contributed unlabelled images, with tissue classes inferred entirely from the generative model.

Prior to atlas learning, all training images were rigidly aligned to MNI space. The Multi-Brain framework jointly estimates a tissue atlas, Gaussian mixture model (GMM) intensity priors, and nonlinear deformation fields, all within a single generative model. The model was configured with $K = 18$ Gaussian components and fit to the pooled training data using default hyper-parameters for the diffeomorphic registration regularisation and bias field

⁵Source code: <https://github.com/WTCN-computational-anatomy-group/mb>

correction, as described in [Brudfors et al., 2020]. After convergence, the 18 learned components were merged by visual inspection into six tissue classes: GM, WM, CSF, bone, soft tissue, and background. Each resulting tissue class is represented by a single spatial prior (one channel of the atlas) but may be modelled by multiple Gaussian components in the intensity domain. For example, the background class is captured by a single spatial prior map but six Gaussians in the intensity model, reflecting the distinct intensity profiles of air, surrounding tissue, and scanner-specific artefacts (Figure 10). This GMM-per-class structure allows each tissue class to capture multimodal intensity distributions without requiring additional spatial prior channels. Full details of the model specification, optimisation procedure, and convergence criteria are provided in [Brudfors, 2020].




Figure 10: Learned Gaussian mixture model intensity priors for CT data. Each colour represents one of the six tissue classes; individual Gaussian components within each class are shown as separate curves of the same colour. For example, the background class (grey) comprises six Gaussian components, reflecting distinct intensity sub-populations (e.g. air, scanner bed).

B.1 Original Groupwise Atlas

The resulting atlas was learned in a data-driven coordinate system that is optimal for the training population but does not correspond to any standard atlas space (Figure 11). A notable feature of this atlas is its extended caudal field of view, which includes the upper cervical spine. This is beneficial for processing clinical CT scans, which frequently extend well below the skull base into the neck, particularly in emergency and oncological settings. The extended coverage ensures that the registration model can anchor to inferior anatomy, improving alignment robustness for scans with variable fields of view and reducing the risk of misregistration when the brain occupies only a portion of the image volume.




Figure 11: Original groupwise CTseg tissue atlas (`mu_CTseg.nii`), learned in a data-driven coordinate system. Note the extended caudal field of view covering the upper cervical spine. Columns: tissue classes (GM, WM, CSF, bone, soft tissue, background). Rows: axial, coronal, and sagittal views.

B.2 MNI-aligned Atlas

For this validation study, the original atlas was nonlinearly aligned to MNI space (Section 4.3). Two MNI-aligned variants are provided:

- `mu_CTseg_spm_mni15.nii`: 1.5 mm isotropic resolution, matching the voxel size of the SPM12 tissue probability atlas. This is the version used throughout this paper.
- `mu_CTseg_spm_mni10.nii`: 1.0 mm isotropic resolution, for applications requiring finer spatial detail, such as cortical morphometry, hippocampal volumetry, or segmentation of thin-slice CT acquisitions where the input resolution exceeds 1.5 mm.

Both MNI-aligned variants crop the atlas to the standard MNI bounding box and therefore do not retain the extended caudal coverage of the original. The combination of the smaller field of view and (for the 1.5 mm version) larger voxels makes the MNI-aligned atlases significantly faster to fit than the original, although the reduced resolution of the 1.5 mm version could in principle lead to less accurate segmentations at fine spatial scales. The 1.0 mm version offers a middle ground: standard-space outputs with higher spatial fidelity, at the cost of longer processing time. The 1.5 mm MNI-aligned atlas used throughout this paper is shown in Figure 12.

All three atlases are distributed with CTseg. The 1.5 mm MNI-aligned version is recommended as the default for most applications; the 1.0 mm version is recommended when finer spatial resolution is needed; and the original groupwise atlas is preferable when extended inferior coverage is required.




Figure 12: MNI-aligned CTseg tissue atlas used in this study. Columns: tissue classes (GM, WM, CSF, bone, soft tissue, background). Rows: axial, coronal, and sagittal views. Brighter voxels indicate higher prior probability.

C Age Prediction

In addition to sex classification (Section 4.4.4), we evaluated age regression using the same GP framework and normalised tissue maps (GP regression with Gaussian likelihood, 10-fold CV). Table 4 and Figure 13 report the results. Age prediction was modest for all three methods, with correlations of $r \approx 0.39$ – 0.45 and RMSE of approximately 10–11 years. All methods performed comparably, and the differences are not meaningful at this sample size ($n = 58$). The limited predictive power is attributable to the narrow, skewed age range of this head-and-neck cancer cohort (predominantly 50–80 years, mean = 64, SD = 12), rather than to the normalisation methods: even the MRI-based pipeline performed similarly. Previous work using the same methodology on larger, age-diverse MRI cohorts [Brudfors et al., 2019] reported substantially stronger age–brain associations. We therefore present these results for completeness but draw no method-specific conclusions from them.

Table 4: Age prediction from normalised tissue maps (10-fold CV, GP regression, FWHM=8 mm).

Method	RMSE (yr)	SD (yr)	Bias (yr)	r
SPM-MR	10.82	10.90	−0.48	0.389
SPM-CT	10.30	10.39	0.04	0.449
CTseg-CT	10.49	10.58	−0.04	0.434




Figure 13: Age prediction: true versus predicted age for each method (10-fold cross-validated GP regression on normalised tissue maps).